

Exploratory Data Analysis on Text Data

Promothesh Chatterjee*

*Copyright© 2024 by Promothesh Chatterjee. All rights reserved. No part of this note may be reproduced, distributed, or transmitted in any form or by any means, including photocopying, recording, or other electronic or mechanical methods, without the prior written permission of the author.

Exploratory Analysis of Text Data

In these notes, we will try to answer some questions based on a hotel reviews dataset.

Assume you are the data analyst for Marriott:

- Identify how consumers view Marriott vis-a-vis competitors Hilton and Hyatt
- Are consumers more happy with competitors than Marriott?
- Do unhappy customers write longer reviews?
- Does it vary by the type of device used to write reviews?
- Can we extract and visualize key words in these reviews?

Let's read the original hotel reviews file in R and take a glimpse of the dataset.

```
# Load up the .CSV data and explore.
library(tidyverse)
hotel_raw <- read_csv("C:/Users/u0474728/Dropbox/Utah Department Stuff/Teaching/Text Analysis/Summer 2017/hotel_data.csv")
glimpse(hotel_raw)
```

```
## Rows: 38,932
## Columns: 5
## $ User_ID      <chr> "id10326", "id10327", "id10328", "id10329", "id10330", "i~
## $ Description  <chr> "The room was kind of clean but had a VERY strong smell o~
## $ Browser_Used <chr> "Edge", "Internet Explorer", "Mozilla", "InternetExplorer~
## $ Device_Used  <chr> "Mobile", "Mobile", "Tablet", "Desktop", "Tablet", "Deskt~
## $ Is Response  <chr> "not happy", "not happy", "not happy", "happy", "not happ~
```

As you can see, there are 38,932 observations (that is 38,932 reviews) and 5 variables all of which have been coded by R as character variables (indicating text variables).

Let's check data to see if there are missing values.

```
sum(!complete.cases(hotel_raw)) # checking the number of incomplete cases
```

```
## [1] 0
```

The negation of function complete cases checks if there is any incomplete cases, and we simply sum it up to find the total number of incomplete cases.

Since there are no missing cases, let's take three relevant variables that are coded as character variables, and let's convert them to Factor (R parlance for categorical variables) as they would be relevant for future analysis.

```

hotel_raw$Is_Response <- as.factor(hotel_raw$Is_Response)
hotel_raw$Device_Used <- as.factor(hotel_raw$Device_Used)
hotel_raw$Browser_Used <- as.factor(hotel_raw$Browser_Used)

```

Let's focus on three hotels (Hilton, Hyatt and Marriott) which are present in the dataset and do a comparative analysis.

```

hotel_raw <- hotel_raw %>%
  mutate(hotel_name = case_when(str_detect(Description, "Hilton")~ "Hilton",
                                str_detect(Description, "Hyatt")~ "Hyatt",
                                str_detect(Description, "Marriott")~ "Marriott"))
hotel_sub <- hotel_raw %>% filter(hotel_name=="Hilton"|hotel_name=="Hyatt"|hotel_name=="Marriott")

```

This R code is manipulating a data frame called “hotel_raw” to create a new data frame called “hotel_sub” that only includes rows where the “hotel_name” column is “Hilton”, “Hyatt”, or “Marriott”.

The code achieves this by first adding a new column called “hotel_name” to the “hotel_raw” data frame using the mutate() function from the dplyr package. The values in this new column are determined using a series of case_when() statements, which test whether certain patterns are found in the “Description” column of the data frame.

For example, the first case_when() statement checks whether the string “Hilton” is found in the “Description” column using the str_detect() function from the stringr package. If this is true, the value “Hilton” is assigned to the “hotel_name” column for that row. Similarly, the second case_when() statement assigns the value “Hyatt” if the string “Hyatt” is found in the “Description” column, and the third case_when() statement assigns the value “Marriott” if the string “Marriott” is found.

Finally, the filter() function from dplyr is used to create the new data frame “hotel_sub” by selecting only the rows from “hotel_raw” where the “hotel_name” column is “Hilton”, “Hyatt”, or “Marriott”.

Now, first, let's take a look at distribution of the class labels (i.e., what proportion of consumers are happy vs. not happy).

```

hotel_sub %>%
  count(Is_Response)%>%
  mutate(freq=n/sum(n))

```

```

## # A tibble: 2 x 3
##   Is_Response     n freq
##   <fct>         <int> <dbl>
## 1 happy         2010 0.618
## 2 not happy     1244 0.382

```

The count function will simply provide a count of variable Is_Response, then we use mutate function from dplyr to create a new variable named ‘freq’ by taking the count and dividing by sum of the count. Roughly, 62% consumers are happy with the hotels and others are not.

Let's see if any hotel is represented disproportionately.

```
hotel_sub %>%
  group_by(hotel_name) %>%
  count(Is_Response)%>%
  mutate(freq=n/sum(n))
```

```
## # A tibble: 6 x 4
## # Groups:   hotel_name [3]
##   hotel_name Is_Response      n freq
##   <chr>      <fct>      <int> <dbl>
## 1 Hilton      happy          847 0.597
## 2 Hilton     not happy          572 0.403
## 3 Hyatt       happy          474 0.629
## 4 Hyatt     not happy          280 0.371
## 5 Marriott   happy          689 0.637
## 6 Marriott   not happy          392 0.363
```

Hilton folks seem to be most unhappy. Before looking at the content of text reviews, let's just examine the distribution of the length of the reviews. We will create a variable called ReviewLength which counts the number of characters in the reviews.

```
hotel_sub<-hotel_sub %>%
  mutate(ReviewLength=nchar(Description))

library(quanteda)
data_corpus_hotelsub <- corpus(hotel_sub, text_field = "Description")# for subsequent analysis
```

Now let's see if the review length varies by the type of device. It could be possible that people write less on a mobile device than on a laptop. We use the function group_by (which is typically used in conjunction with summarise). Using function summarize, we create a variable called AvLength which is just a mean of ReviewLength.

```
hotel_sub%>%
  group_by( hotel_name, Device_Used) %>%
  summarise(AvLength=mean(ReviewLength))
```

```
## # A tibble: 9 x 3
## # Groups:   hotel_name [3]
##   hotel_name Device_Used AvLength
##   <chr>      <fct>      <dbl>
## 1 Hilton      Desktop          1201.
## 2 Hilton      Mobile           1189.
```

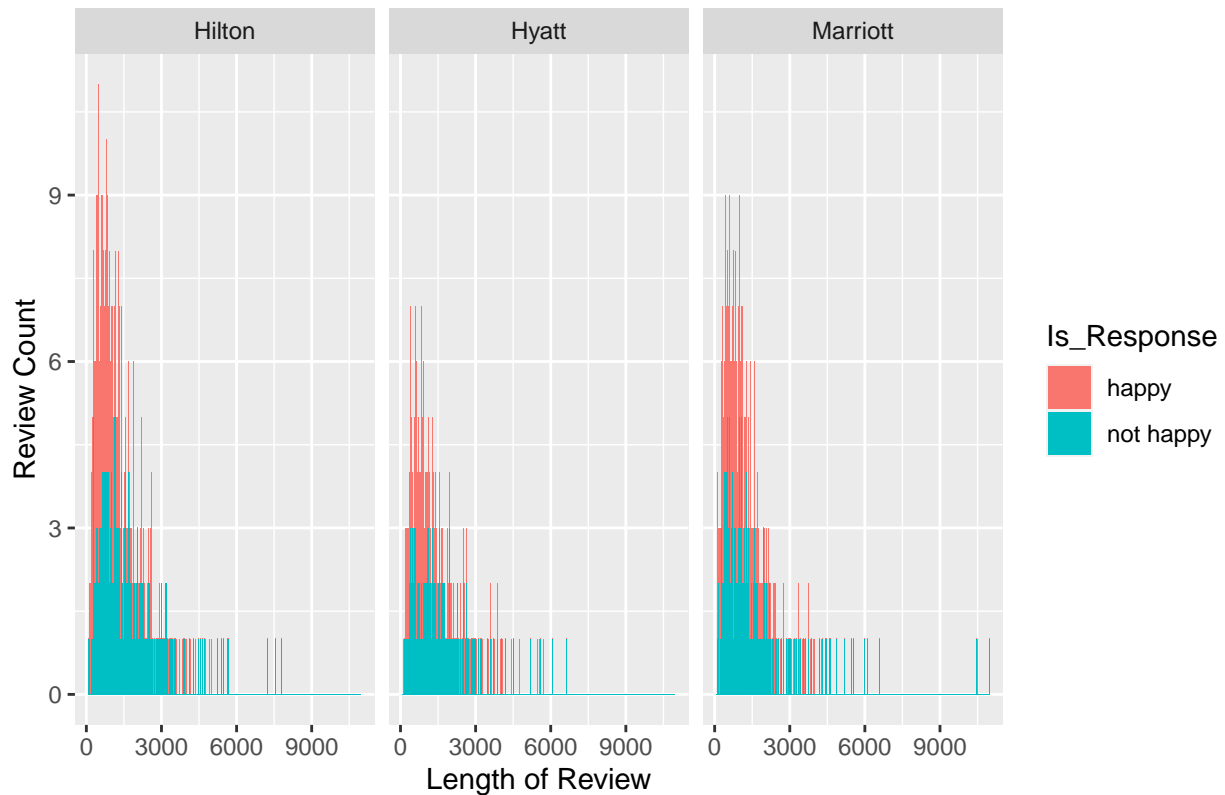
```
## 3 Hilton      Tablet      1241.  
## 4 Hyatt       Desktop     1268.  
## 5 Hyatt       Mobile      1213.  
## 6 Hyatt       Tablet      1236.  
## 7 Marriott   Desktop     1113.  
## 8 Marriott   Mobile      1195.  
## 9 Marriott   Tablet      1209.
```

The output suggests that there is a minimal difference between number of characters in each review between the hotels and devices.

Now let's check if review length differs based on whether consumers are happy or not with the hotel. We might suspect that people who are not happy might post longer reviews, let's visualize that using the tidyverse package GGPlot. The first argument of ggplot is the data frame `hotel_raw`, next we indicate the x variable in the function aesthetics (`aes`) and use the command `fill` to separate the plot by the factor variable `Is_Response`.

```
library(ggplot2)  
  
ggplot(hotel_sub, aes(x = ReviewLength, fill = Is_Response)) +  
  geom_histogram(binwidth = 5) + facet_wrap(~ hotel_name) +  
  labs(y = "Review Count", x = "Length of Review",  
       title = "Distribution of Review Lengths with Class Labels")
```

Distribution of Review Lengths with Class Labels



The relationship between length of reviews and customer happiness does not seem to be clear cut. While the number of happy customers is definitely more than not-happy customers, the length of the review of not-happy customers seems to be slightly more only in the case of Marriott.

So far, we haven't even gone to NLP. Let's now do the preprocessing (we have seen these in the previous section). We essentially tokenize, remove stopwords and do other preprocessing to create document-feature-matrix.

```
library(quantda)
hotel_sub_tokens <- quantda::tokens(hotel_sub$Description, what = "word",
                                   remove_numbers = TRUE, remove_punct = TRUE,
                                   remove_symbols = TRUE)

#Create DFM
hotel_sub_dfm <- hotel_sub_tokens %>%
  tokens_remove(stopwords(source = "smart")) %>%
  tokens_wordstem() %>%
  tokens_tolower() %>%
  dfm()

dim(hotel_sub_dfm)
```

```
## [1] 3254 14158
```

```
head(hotel_sub_dfm, n = 5)
```

```
## Document-feature matrix of: 5 documents, 14,158 features (99.35% sparse) and 0 docvars.
##           features
## docs   wonder staff great locat defin price high standard hotel free
## text1      1     1     2     1     1     1     1     1     1     2     2
## text2      0     0     0     0     0     2     0     0     0     9     4
## text3      0     3     0     0     0     0     1     0     0     0     0
## text4      0     0     0     0     0     1     0     0     0     4     0
## text5      0     1     1     1     0     0     2     0     0     4     2
## [ reached max_nfeat ... 14,148 more features ]
```

Check the presence of the word ‘hotel’ multiple times across reviews (refer to the notes on TFIDF, we will perform TFIDF shortly).

Concordances Concordancing in Text Analysis involves extracting words from one or more texts (Lindquist 2009). Typically, these concordances are presented as Key Word In Context (KWIC), which displays the search term along with some preceding and following context. As a result, these displays are commonly known as Key Word In Context (KWIC) concordances.

Concordancing allows one to observe the usage of a term in the data, examine the frequency of a particular word in a text or group of texts, and extract examples. It also serves as a fundamental procedure and frequently serves as the initial step in more advanced analyses of language data.

```
hilt <- data_corpus_hotelsub %>%
  kwic(pattern = "Hilton.*", window = 5, valuetype = "regex")
```

You can check out the hilt dataset to understand the context in which the word Hilton appears.

We can also search for phrase patterns.

```
kwic_hiltonhonors <- quanteda::kwic(x = data_corpus_hotelsub,
  pattern = quanteda::phrase("Hilton Honors"),
  window = 5)
```

Word Frequency The majority of techniques employed in text analysis heavily rely on frequency data. Consequently, identifying the most common words present in a given text is a fundamental approach in text analytics. Essentially, frequency data constitutes the backbone of text analysis, frequently appearing in the form of word frequency lists that consist of word forms and their corresponding frequency in a given text or text collection.

```
dfmat2 <- corpus_subset(data_corpus_hotelsub, hotel_name == "Hilton") %>%
tokens(remove_punct = TRUE) %>%
tokens_remove(stopwords("en")) %>%
dfm()

library(quantda.textstats)
tstat1 <- textstat_frequency(dfmat2, groups =hotel_name)
head(tstat1, 10)
```

```
##      feature frequency rank docfreq  group
## 1    hotel      3206     1    1128 Hilton
## 2     room      2954     2    1030 Hilton
## 3   hilton      2010     3    1358 Hilton
## 4      --      1397     4     603 Hilton
## 5     stay      1238     5     794 Hilton
## 6      one      1004     6     589 Hilton
## 7    great       960     7     581 Hilton
## 8    staff       932     8     680 Hilton
## 9     nice       869     9     532 Hilton
## 10   rooms       861    10     597 Hilton
```

This code snippet performs the following tasks:

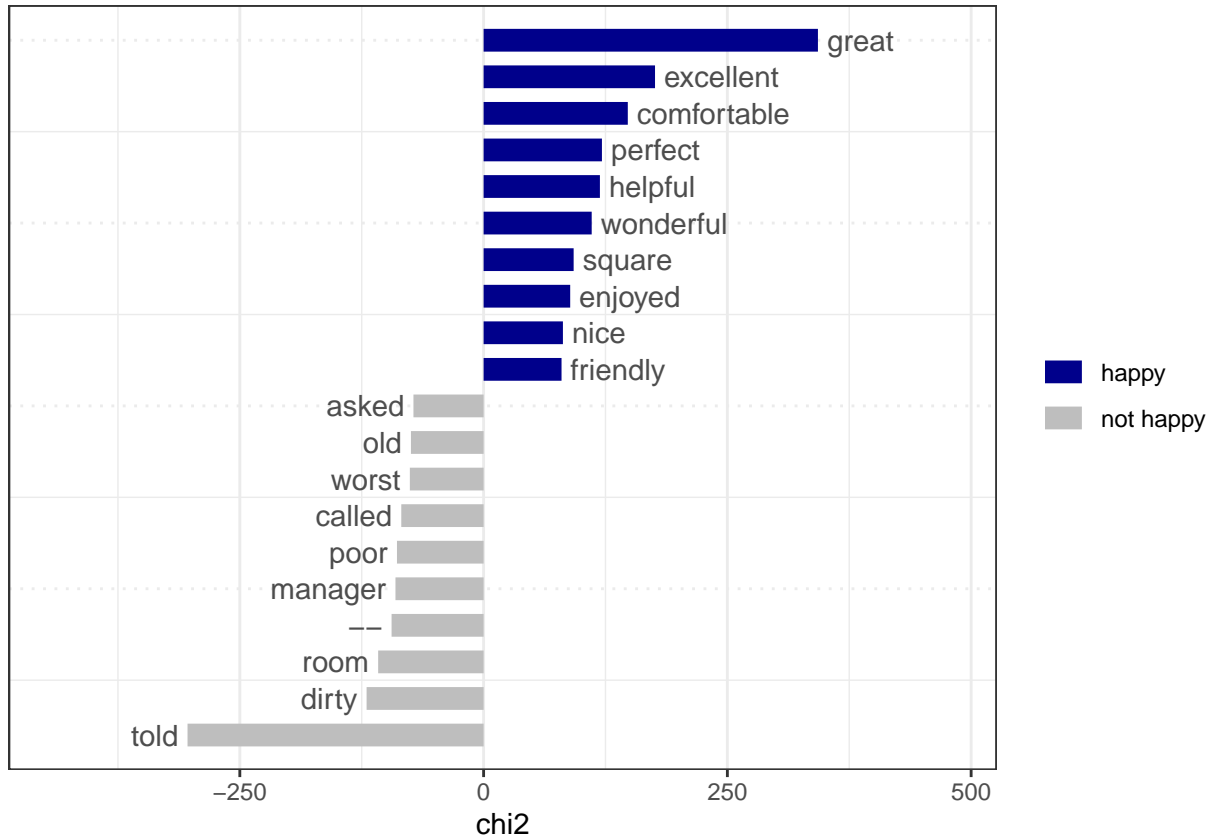
1. `corpus_subset(data_corpus_hotelsub, hotel_name == "Hilton")` filters a corpus `data_corpus_hotelsub` to only include documents where the `hotel_name` metadata variable is equal to “Hilton”. This creates a new corpus object `dfmat2`.
2. `%>%` is a pipe operator that passes the output of the previous function to the next function.
3. `tokens(remove_punct = TRUE)` converts the corpus into a token object and removes all punctuation marks.
4. `tokens_remove(stopwords("en"))` removes all stopwords (common words that are not very informative) using the English language stopwords list.
5. `dfm()` converts the token object to a document-feature matrix (dfm), which is a matrix where the rows represent the documents and the columns represent the features (words in this case).
6. `textstat_frequency(dfmat2, groups = hotel_name)` calculates the frequency of each word in the `dfmat2` matrix and groups them by `hotel_name`. This creates a new object `tstat1` which contains the frequency counts for each word grouped by the hotel name.
7. `head(tstat1, 10)` displays the top 10 most frequent words for each hotel in the `tstat1` object.

Word Cloud Let's now generate a word cloud for further visualization. The idea is pretty simple, the word size indicates the frequency of occurrence of the word. We will use the `textplot_wordcloud` function from `quanteda` (again, you check the help file by typing `?textplot_wordcloud` in the console to look at the details of the arguments passed)

```
library(quanteda.textplots)
textplot_wordcloud(
  hotel_sub_dfm,
  min_size = 0.5,
  max_size = 4,
  min_count = 3,
  max_words = 200,
  color = "darkblue",
  font = NULL,
  adjust = 0,
  rotation = 0.1,
  random_order = FALSE,
  random_color = FALSE,
  ordered_color = FALSE,
  labelcolor = "gray20",
  labelsizes = 1.5,
  labeloffset = 0,
  fixed_aspect = TRUE,
  comparison = FALSE
)
```



```
library(quantda.textstats)
tstat1 <- textstat_keyness(dfmat2, target = "happy")
textplot_keyness(tstat1, margin = 0.2, n = 10)
```



Let's get the frequency weights

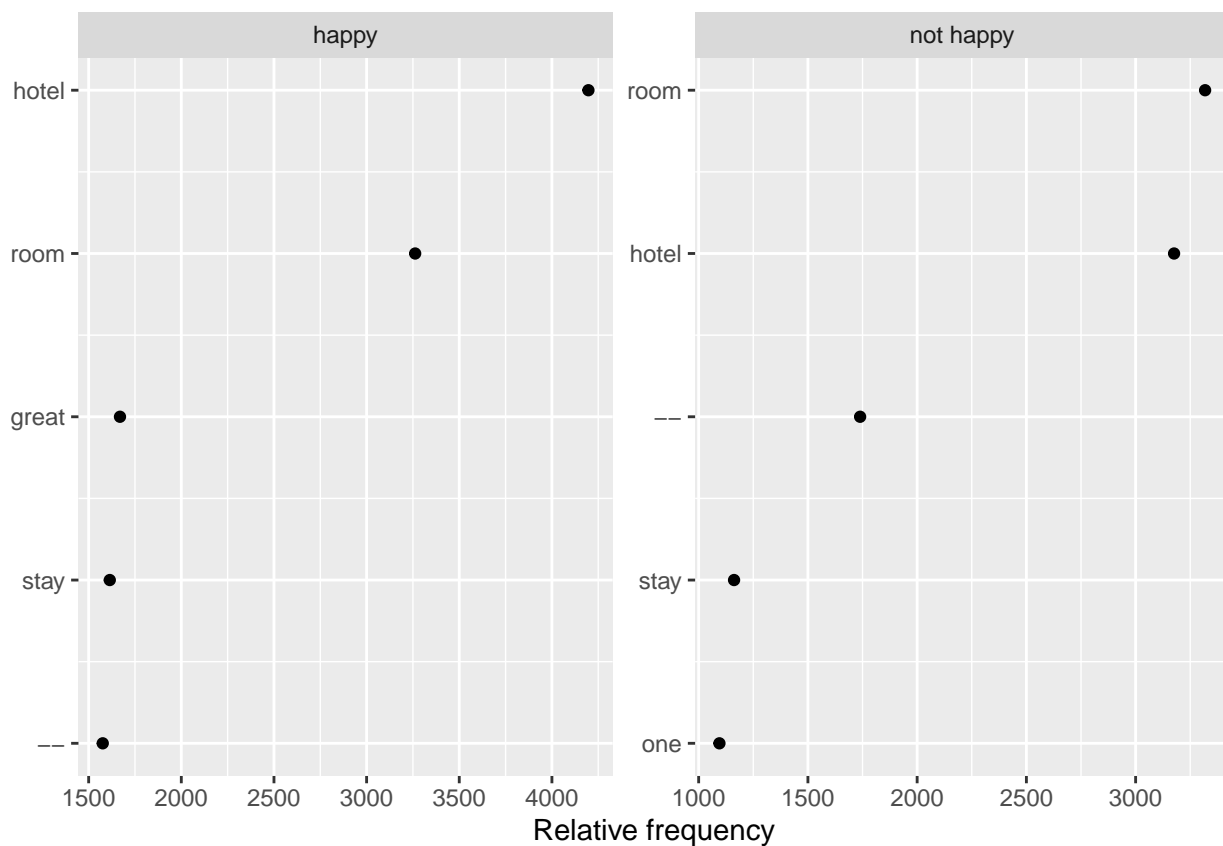
```
freq_weight <- quantda.textstats::textstat_frequency(dfmat2,
  n = 5,
  groups = dfmat2$Is_Response)
freq_weight
```

##	feature	frequency	rank	docfreq	group
## 1	hotel	4197	1	1	happy
## 2	room	3262	2	1	happy
## 3	great	1669	3	1	happy
## 4	stay	1614	4	1	happy
## 5	--	1576	5	1	happy
## 6	room	3318	1	1	not happy
## 7	hotel	3176	2	1	not happy
## 8	--	1739	3	1	not happy
## 9	stay	1162	4	1	not happy

```
## 10      one      1095    5      1 not happy
```

GGPlot these

```
ggplot(freq_weight, aes(nrow(freq_weight):1, frequency)) +  
  geom_point() +  
  facet_wrap(~ group, scales = "free") +  
  coord_flip() +  
  scale_x_continuous(breaks = nrow(freq_weight):1,  
                     labels = freq_weight$feature) +  
  labs(x = NULL, y = "Relative frequency")
```



This code is an example of how to create a scatterplot of relative frequencies for a set of features, grouped by a categorical variable, using the ggplot2 package in R.

Here's a breakdown of the code:

1. `ggplot(freq_weight, aes(nrow(freq_weight):1, frequency))`: This sets up the ggplot object and specifies the dataset (`freq_weight`) and the x and y variables for the plot. The `nrow(freq_weight):1` part creates a sequence of numbers from the number of rows in the dataset down to 1, which will be used as the x-axis values. The `frequency` variable is used as the y-axis values.
2. `geom_point()`: This adds a layer to the plot, which will display the data as points.

3. `facet_wrap(~ group, scales = "free")`: This creates a separate panel in the plot for each unique value of the `group` variable in the dataset, and sets the scales to be “free” (i.e., they will be scaled independently for each panel).
4. `coord_flip()`: This flips the x and y axes, so that the plot is displayed horizontally.
5. `scale_x_continuous(breaks = nrow(freq_weight):1, labels = freq_weight$feature)`: This sets the breaks and labels for the x-axis. The `breaks` argument specifies the location of tick marks on the x-axis (which will be the sequence of numbers created earlier), and the `labels` argument specifies the labels for each tick mark, which will be the `feature` variable in the `freq_weight` dataset.
6. `labs(x = NULL, y = "Relative frequency")`: This sets the axis labels, with `x = NULL` indicating that the x-axis should not be labeled, and `y = "Relative frequency"` indicating that the y-axis should be labeled “Relative frequency”.

Overall, this code will produce a scatterplot of relative frequencies for each feature, grouped by the value of a categorical variable, with separate panels for each value of the categorical variable. The x-axis will display the feature names, and the y-axis will display the relative frequency values.

N-grams and Collocations N-grams, collocations, and keyness are all related concepts in linguistics. Collocation pertains to the occurrence of words together, such as the commonly used phrase “touch base”. This phrase is an example of a collocation because the words “touch” and “base” frequently appear together, more often than by chance alone.

In contrast, n-grams represent a broader set of word combinations that occur together. For instance, bi-grams refer to two words that are frequently used together, while tri-grams refer to three words that are commonly used together, and so on.

```
hotel_ngrams<-toks_hotel %>%
  tokens_ngrams(n=2)
hotel_ngrams_dfm<-dfm(hotel_ngrams)
tstat1 <- textstat_frequency(hotel_ngrams_dfm)
head(tstat1, 10)
```

```
##           feature frequency rank docfreq group
## 1           $_--      1003     1      611   all
## 2       front_desk       882     2      681   all
## 3           $_---       393     3       285   all
## 4       room_service       335     4       267   all
## 5           --_--       324     5       234   all
## 6       --_minutes       299     6       244   all
## 7 walking_distance       287     7       269   all
## 8       across_street       284     8       256   all
## 9       times_square       261     9       183   all
## 10      new_york         256    10       169   all
```

We can get a similar output even with general collocations code

```
text_coll <- textstat_collocations(toks_hotel, size = 2, min_count = 20)
text_coll %>% arrange(across(starts_with("count"), desc)) %>% head()
```

##	collocation	count	count_nested	length	lambda	z
## 1	front desk	881	0	2	7.939824	90.66670
## 30	room service	331	0	2	2.524822	40.58670
## 11	walking distance	287	0	2	9.371667	51.57749
## 3	across street	284	0	2	7.053361	65.00979
## 2	times square	259	0	2	7.045201	66.23640
## 396	new york	253	0	2	11.138878	17.48291

We can also run a customized ngram

```
cust_bigram<- tokens_compound(hotel_ngrams, phrase("room*"))
cust_bigram<- tokens_select(cust_bigram, phrase("room_*"))

cust_bigram<-dfm(cust_bigram)
tstat1 <- textstat_frequency(cust_bigram)
head(tstat1, 10)
```

##	feature	frequency	rank	docfreq	group
## 1	room_service	335	1	267	all
## 2	room_clean	182	2	181	all
## 3	room_nice	131	3	127	all
## 4	room_ready	85	4	72	all
## 5	room_--th	69	5	68	all
## 6	room_spacious	66	6	66	all
## 7	room_room	66	6	63	all
## 8	room_small	60	8	59	all
## 9	room_great	58	9	58	all
## 10	room_large	55	10	54	all

```
# You can visualize it as a word cloud as well
# textplot_wordcloud(cust_bigram, min_size = 0.5, max_size = 4,
# min_count = 3, max_words = 200, color = "darkblue", font = NULL,
# adjust = 0, rotation = 0.1, random_order = FALSE, random_color = FALSE,
# ordered_color = FALSE, labelcolor = "gray20", labelsizes = 1.5,
# labeloffset = 0, fixed_aspect = TRUE, comparison = FALSE)
```

We have seen that even without getting into any fancy analysis, we can get a basic understanding from exploratory data analysis.